

Epsilon-Greedy Exploration: Agent selects action: $\arg \max Q_{eval}$ with prob. $(1 - \epsilon)$, random with prob. ϵ



Experience Replay: Store transitions (S, A, R, S') in buffer $\mathcal{M}$ and sample mini-batches



Target Network Update: Periodically copy $\theta^{eval} \rightarrow \theta^{target}$ for stable learning